

Raising Polynomials to Powers

Alex Pan

April 2024

The problem

To determine properties about some invariant I don't understand yet, we need to compute the function: Given a polynomial $f = \mathbb{F}_p[x_1, x_2, \dots, x_n]$ with homogeneous degree n , and a prime p , we need to compute $g = f^{p-1}$, lift g to $\mathbb{Z}[x_1, x_2, \dots, x_n]$, and compute g^p .

Method 1: Jack's C code

Don't know much about this but it does the entire worst case in approximately 800 milliseconds on his machine, so this is already pretty fast. It is also fast enough already when the number of starting terms is small.

Method 2: Multinomial Theorem

This method involves precomputing the multinomial coefficients for a certain number of terms. However, precomputing the multinomial coefficients isn't very parallelizable, as well as slow. Even for the lowest case of interest, being 4 variables and a prime of 5, we need to store $35 * \binom{20+35-1}{35-1} \approx 1.1 * 10^6$ integers, which corresponds to about 10 petabytes of storage. So, this is not the most viable solution.

Though, I will mention that with pregenerated coefficients, multinomial theorem is by far the fastest method for the worst case of the first step $g = f^{p-1}$ at approximately 79 microseconds, though FFT does it in 1 millisecond anyways.

Method 3: Repeated Squaring with GPU-parallelized Trivial Multiplication

This just sucks, maybe its comparable when on a research cluster, since the results for all the steps on my machine are an order of magnitude behind the results for the same steps in the papers I've read about this method

Using FFT to raise polynomials to powers

DFT works for multiplying polynomials by evaluating both polynomials at n points, represented by the output vectors \hat{f} and \hat{g} . Component-wise multiplication is then performed on \hat{f} and \hat{g} , representing the product polynomial evaluated at n points, and IDFT is performed to give the actual coefficients of the polynomial.

If we want to raise a polynomial to a power, we can simply DFT to get \hat{f} , then raise each element of \hat{f} to the p -th power, and that will represent f^p evaluated at n points, and IDFT to get the coefficients of f^p .

This should be much faster than the repeated squaring method; this method only requires one DFT to be evaluated, and raising each component to a power in parallel should be effectively instant. In contrast to $\log_2(p)$ different DFTs

Proofs for correctness of computed numbers:

Because the GPU doesn't report when an integer datatype overflows, we need to verify that the numbers the algorithm returns are actually correct, and less than 2^{64} .

Theorem: Let f be a k -variate polynomial with homogeneous degree n . Let M be the coefficient of f^p . Then, the maximum coefficient of f^p satisfies $M' \leq M^p \cdot \binom{n+k-1}{k-1}^{p-1}$

Proof: Say we have a k -variate polynomial f with homogeneous degree n . Each term of f will look like $cx_1^{d_1}x_2^{d_2}\dots x_k^{d_k}$, with $\sum_1^k d_i = n$. To make things easier, let (d_1, d_2, \dots, d_k) denote the degree of a term from now on.

Let $(d'_1, d'_2, \dots, d'_k)$ denote the degree of a term of f^2 . To find all unreduced terms that contribute to it, we look at pairs of terms of f with the original degrees $(d'_1 - a_1, d'_2 - a_2, d'_3 - a_3, \dots, d'_k - a_k)$ and $(a_1, a_2, a_3, \dots, a_k)$, where $a_i > 0$ and $\sum_1^k a_i = n$. The number of these terms is bounded above by the number of weak integer compositions of n into k parts, which is given by $\binom{n+k-1}{k-1}$. The coefficients of each of these resulting terms is bounded above by M^2 , and because the number of these is bounded above by $\binom{n+k-1}{k-1}$, we know that the maximum reduced coefficient of f^2 is $M^2 \cdot \binom{n+k-1}{k-1}$

The inductive step is similar. Let $(d'_1, d'_2, \dots, d'_k)$ denote the degree of a term of f^p . The pairs of terms that contribute to it come respectively from f^{p-1} and f , so we look again at a term of f^{p-1} with degree $(d'_1 - a_1, d'_2 - a_2, d'_3 - a_3, \dots, d'_k - a_k)$, and at a term of f with degree $(a_1, a_2, a_3, \dots, a_k)$. Again, there are at most $\binom{n+k-1}{k-1}$ of these, and because the maximum unreduced coefficient is bounded above by $(M^{p-1} \cdot \binom{n+k-1}{k-1}^{p-2})$.

(M), we have an upper bound of $M^p \cdot \binom{n+k-1}{k-1}^{p-1}$ for the coefficients. QED

We are interested in two steps: raising $g = f^{p-1} \pmod p$ for primes p , and raising g^p for the same p .

- $g = f^{p-1} \pmod p$

If f is a polynomial of 4 variables, then we have an upper bound of $(p-1)^{p-1} \cdot \binom{4+p-1}{4-1}^{p-1}$. This is less than 2^{64} for primes 5 and 7, so we can reduce mod p at the end of our computation without overflowing. (In the raise to $p-1$ case, the OEIS sequence gives a better bound, but I don't know how to prove that the problems are equivalent yet. That bound is actually obtainable)

If f is a polynomial of 5 variables, then primes 5 and 7 still work for this first step.

- g^p

If f is a polynomial of 4 variables, and it has been raised to the $p-1$ and taken mod p , then the greatest coefficient of g is $p-1$, and g is homogeneous with degree $4 * p$. This gives the upper bound $(p-1)^p \cdot \binom{4p+4-1}{4-1}^p$. For $p = 5$, this evaluates to $1.00733564 \cdot 10^{17}$, which is less than $2^{64} - 1$, or $1.84467441 \cdot 10^{19}$. So, for the case of 4 variables, and $p = 5$, Int64's can be used without having to worry about overflowing.

Update 6/24: Instead of all of that, I could've just calculated the maximum coefficient using some programming

Table of bounds for first step